

Assignment 2 – Report

key-value pairs that were sent from the mappers to the reducers:

step 1:

number of pairs – 212796585

size (bytes) - 4781939956

step2:

number of pairs – 4125612

size (bytes) – 238193533

step3:

number of pairs – 1181055

size (bytes) – 69416660

step4:

number of pairs – 334360

size (bytes) - 12730612

Scalability:

1. Running with 7 instances: 9:44 minutes.
2. Running with 3 instances: 24:21 minutes.

Interesting word pairs: